

Introduction to R*

Part 3: Matrices/Arrays & Special Data Types

Wim R.M. Cardoen

Last updated: 07/22/2024 @ 11:48:36

Contents

1	Matrices & Arrays	2
1.1	Creation of matrices	2
1.1.1	Examples	2
1.1.2	Matrices: vectors with a non-NULL dim attribute	3
1.2	Retrieving elements/subsetting	5
1.2.1	Examples	5
1.2.2	Exercises	9
1.3	Operations on matrices	9
1.3.1	Examples	9
1.3.2	Exercises	12
1.4	Hash tables/dictionaries	14
1.4.1	Examples	14
1.5	Arrays	15
2	Special Data Types (Factors and Date/Time types)	16
2.1	Attributes	16
2.1.1	Examples	16
2.2	Factor variables (Categorical variables)	19
2.2.1	Examples	19
2.3	Dates and times in R	21
2.3.1	Examples	21
	Bibliography	23

*© - Wim R.M. Cardoen, 2022 - The content can neither be copied nor distributed without the **explicit** permission of the author.

"It is my experience that proofs involving matrices can be shortened by 50% if one throws the matrices out" (Emil Artin)

1 Matrices & Arrays

Matrices and arrays are **homogeneous atomic vectors** with an **extra** attribute: dimension

By default, the elements are stored in a **column-major** fashion. (cfr. **Fortran**). However, we can store the elements in **row-major** order (cfr. **C**) as well.

1.1 Creation of matrices

Matrices can be created in different ways:

- use of the `matrix()` function
- use of `rbind()/cbind()`
- set the `attributes()` of a vector
- special functions like e.g. `diag()`

1.1.1 Examples

- use of the `matrix()` function:

The `matrix()` function creates a matrix based on a vector.

By default, the elements are stored in a **column-major** fashion.

The use of the flag `byrow=TRUE` will store the data in a **row-major** fashion.

```
A <- matrix(data=1:10, nrow=2)    # Column-major (like Fortran)
A
```

```
      [,1] [,2] [,3] [,4] [,5]
[1,]    1    3    5    7    9
[2,]    2    4    6    8   10
```

```
B <- matrix(data=c(2,3,893,0.17), nrow=2, ncol=2)
B
```

```
      [,1] [,2]
[1,]    2 893.00
[2,]    3  0.17
```

```
C <- matrix(data=1:10, nrow=2, byrow=TRUE)    # Row-major (like C, C++)
C
```

```
      [,1] [,2] [,3] [,4] [,5]
[1,]    1    2    3    4    5
[2,]    6    7    8    9   10
```

- use of the `rbind()/cbind()` functions:
 - `rbind()`: Bind several vectors (as rows) into a matrix.
 - `cbind()`: Bind several vectors (as columns) into a matrix.

```
A <- rbind(1:10,11:20)
A
```

```
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]
[1,]     1     2     3     4     5     6     7     8     9     10
[2,]    11    12    13    14    15    16    17    18    19    20
```

```
typeof(A)
```

```
[1] "integer"
```

```
class(A)
```

```
[1] "matrix" "array"
```

```
B <- cbind(1:5,6:10,11:15)
B
```

```
      [,1] [,2] [,3]
[1,]     1     6    11
[2,]     2     7    12
[3,]     3     8    13
[4,]     4     9    14
[5,]     5    10    15
```

```
class(B)
```

```
[1] "matrix" "array"
```

1.1.2 Matrices: vectors with a non-NULL dim attribute

The **fundamental** difference between an R vector and matrix is the presence (in the case of matrices) of a non NULL **dim** attribute.

We can easily convert a vector into a matrix by setting the dimensions of the vector:

- through the **dim()** function.
- through the **attr()** function.

The inverse can be done as well by setting the **dim** attribute of matrix to NULL.

```
A <- 1:10
typeof(A)
```

```
[1] "integer"
```

```
class(A)
```

```
[1] "integer"
```

```
dim(A)
```

```
NULL
```

```
# Matrix
B <- matrix(1:10,nrow=2,ncol=5,byrow=TRUE)
typeof(B)
```

```
[1] "integer"
```

```
class(B)
```

```
[1] "matrix" "array"
```

```
dim(B)
```

```
[1] 2 5
```

```
# Vector
```

```
A <- 1:10
```

```
A
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

```
dim(A)
```

```
NULL
```

```
typeof(A)
```

```
[1] "integer"
```

```
class(A)
```

```
[1] "integer"
```

```
# OPTION I: Using the dim function transform a vector into a matrix
```

```
dim(A) <- c(2,5)
```

```
A
```

```
      [,1] [,2] [,3] [,4] [,5]  
[1,]    1    3    5    7    9  
[2,]    2    4    6    8   10
```

```
dim(A)
```

```
[1] 2 5
```

```
typeof(A)
```

```
[1] "integer"
```

```
class(A)
```

```
[1] "matrix" "array"
```

```
# Converting the matrix back to a vector
```

```
dim(A) <- NULL
```

```
dim(A)
```

```
NULL
```

```
typeof(A)
```

```
[1] "integer"
```

```
class(A)
```

```
[1] "integer"
```

```
# Option II: More general way
```

```
# Convert vector into a matrix
```

```
A <- 1:8
```

```
A
```

```
[1] 1 2 3 4 5 6 7 8
```

```
class(A)
```

```
[1] "integer"
```

```
attr(A, 'dim') <- c(2,4)
```

```
A
```

```
      [,1] [,2] [,3] [,4]  
[1,]     1     3     5     7  
[2,]     2     4     6     8
```

```
class(A)
```

```
[1] "matrix" "array"
```

```
# Convert matrix into a vector.
```

```
attr(A, 'dim') <- NULL
```

```
A
```

```
[1] 1 2 3 4 5 6 7 8
```

```
class(A)
```

```
[1] "integer"
```

1.2 Retrieving elements/subsetting

Matrices (and arrays) can be subsetted in different ways:

- use an **index** for each dimension, where the dimensions are comma-separated
 - If an **index** for a dimension is **omitted**:
consider all dimensions (may lead to reduction of the dimension)
 - **but** you can use **drop=FALSE** to prevent dimensionality reduction.
- use another **vector** (can be either linear or a vector for each dimension)
- by using another **matrix**.

1.2.1 Examples

- Use of **indices**:

```
A <- matrix(1:30, nrow=6, ncol=5)
A
```

```
      [,1] [,2] [,3] [,4] [,5]
[1,]     1     7    13    19    25
[2,]     2     8    14    20    26
[3,]     3     9    15    21    27
[4,]     4    10    16    22    28
[5,]     5    11    17    23    29
[6,]     6    12    18    24    30
```

```
A[3,4]
```

```
[1] 21
```

```
A[6,2]
```

```
[1] 12
```

```
x1 <- A[2,]
x1
```

```
[1]  2  8 14 20 26
```

```
dim(x1)
```

```
NULL
```

```
x2 <- A[,3]
x2
```

```
[1] 13 14 15 16 17 18
```

```
dim(x2)
```

```
NULL
```

The flag **drop=FALSE** can be used to prevent dimensionality reduction

```
y1 <- A[2,,drop=FALSE]
y1
```

```
      [,1] [,2] [,3] [,4] [,5]
[1,]     2     8    14    20    26
```

```
dim(y1)
```

```
[1] 1 5
```

```
y2 <- A[,3,drop=FALSE]
y2
```

```
      [,1]
[1,]    13
[2,]    14
[3,]    15
[4,]    16
[5,]    17
[6,]    18
```

```
dim(y2)
```

```
[1] 6 1
```

- Use of **vector(s)**:

```
A
```

```
      [,1] [,2] [,3] [,4] [,5]
[1,]     1     7    13    19    25
[2,]     2     8    14    20    26
[3,]     3     9    15    21    27
[4,]     4    10    16    22    28
[5,]     5    11    17    23    29
[6,]     6    12    18    24    30
```

```
x1 <- A[2:4,]
x1
```

```
      [,1] [,2] [,3] [,4] [,5]
[1,]     2     8    14    20    26
[2,]     3     9    15    21    27
[3,]     4    10    16    22    28
```

```
dim(x1)
```

```
[1] 3 5
```

```
x2 <- A[,1:3]
x2
```

```
      [,1] [,2] [,3]
[1,]     1     7    13
[2,]     2     8    14
[3,]     3     9    15
[4,]     4    10    16
[5,]     5    11    17
[6,]     6    12    18
```

```
dim(x2)
```

```
[1] 6 3
```

```
# Using a vector for EACH dimension
A[c(1,3),c(2,4)]
```

```
      [,1] [,2]
[1,]    7   19
[2,]    9   21
```

```
# Using 1 vector => Linear index
A[c(1,3,8,10)]
```

```
[1]  1  3  8 10
```

```
A[c(TRUE,FALSE,TRUE,TRUE,FALSE,TRUE),c(2,3)]
```

```
      [,1] [,2]
[1,]    7   13
[2,]    9   15
[3,]   10   16
[4,]   12   18
```

```
A[c(TRUE,FALSE,TRUE,TRUE,FALSE,TRUE),]
```

```
      [,1] [,2] [,3] [,4] [,5]
[1,]    1    7   13   19   25
[2,]    3    9   15   21   27
[3,]    4   10   16   22   28
[4,]    6   12   18   24   30
```

```
# Use of a linear index
A[c(TRUE,FALSE,TRUE,TRUE,FALSE,TRUE)]
```

```
[1]  1  3  4  6  7  9 10 12 13 15 16 18 19 21 22 24 25 27 28 30
```

- Use of a **matrix**:

```
A
```

```
      [,1] [,2] [,3] [,4] [,5]
[1,]    1    7   13   19   25
[2,]    2    8   14   20   26
[3,]    3    9   15   21   27
[4,]    4   10   16   22   28
[5,]    5   11   17   23   29
[6,]    6   12   18   24   30
```



```
mysubset <- matrix(c( 2, 1,
                     3, 5,
                     4, 2,
                     6, 5), ncol=2, byrow=TRUE )
A[mysubset]
```

```
[1] 2 27 10 30
```

1.2.2 Exercises

- Create the following matrix A, given by:

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	3	9	27	81	243	729
[2,]	5	25	125	625	3125	15625
[3,]	7	49	343	2401	16807	117649
[4,]	11	121	1331	14641	161051	1771561
[5,]	13	169	2197	28561	371293	4826809
[6,]	17	289	4913	83521	1419857	24137569

1. get element 343
2. get the elements 25, 625, 2197 and 4826809 (all at once).
3. get the fourth row as a vector.
4. get the fourth row as a matrix.
5. get columns 2 and 3 (at the same time).
6. get everything except rows 2 and 4.
7. the diagonal of matrix A.

1.3 Operations on matrices

- Operations like `*`, `/`, `+` happen element-wise.
- There are also more specialized functions:
 - the mean over rows and columns (`rowMeans()`, `colMeans()`)
 - linear algebra functions (`%*%`, `t()`, ...)

1.3.1 Examples

- Operations (by **default: element-by-element**):

```
A <- matrix(1:10, nrow=2)
B <- matrix( seq(10, 100, by=10), nrow=2)
A
```

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	1	3	5	7	9
[2,]	2	4	6	8	10

```
B
```

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	10	30	50	70	90
[2,]	20	40	60	80	100

```
A*B
```

```
      [,1] [,2] [,3] [,4] [,5]  
[1,]   10   90  250  490  810  
[2,]   40  160  360  640 1000
```

```
C <- matrix(rep(2,10), nrow=2)  
C
```

```
      [,1] [,2] [,3] [,4] [,5]  
[1,]     2     2     2     2     2  
[2,]     2     2     2     2     2
```

```
C**A
```

```
      [,1] [,2] [,3] [,4] [,5]  
[1,]     2     8    32   128   512  
[2,]     4    16    64   256  1024
```

- Calculate row and column means :

```
# Means of rows and columns
```

```
A
```

```
      [,1] [,2] [,3] [,4] [,5]  
[1,]     1     3     5     7     9  
[2,]     2     4     6     8    10
```

```
rowMeans(A)
```

```
[1] 5 6
```

```
colMeans(A)
```

```
[1] 1.5 3.5 5.5 7.5 9.5
```

- Matrix multiplication (%*%) :

```
A <- matrix(1:6, nrow=2)
```

```
A
```

```
      [,1] [,2] [,3]  
[1,]     1     3     5  
[2,]     2     4     6
```

```
B <- matrix(seq(10,120,by=10), nrow=3)
```

```
B
```

```
      [,1] [,2] [,3] [,4]  
[1,]   10   40   70  100  
[2,]   20   50   80  110  
[3,]   30   60   90  120
```

```
C <- A%*%B
```

```
C
```

```
      [,1] [,2] [,3] [,4]  
[1,]  220  490  760 1030  
[2,]  280  640 1000 1360
```

```
dim(C)
```

```
[1] 2 4
```

- **Linear algebra** routines

Some of the more common ones in R:

- `solve()` : **invert** a square matrix
- `diag()`
 - **extracts** the diagonal of a matrix when a matrix is provided.
 - **creates** a diagonal matrix when a vector is provided.
- `eigen()` : calculates the **eigenvalues** and **eigenvectors** of a matrix
- `det()` : calculates the **determinant** of a matrix.
- `t()`: calculates the **transpose**¹ of a matrix.

```
# Invert matrix A
A <- matrix(c(1, 3, 2, 4), ncol = 2, byrow = T)
Ainv <- solve(A)
Ainv %*% A
```

```
      [,1] [,2]
[1,]     1     0
[2,]     0     1
```

```
# Create a diagonal matrix
C <- diag(c(1,4,7))
C
```

```
      [,1] [,2] [,3]
[1,]     1     0     0
[2,]     0     4     0
[3,]     0     0     7
```

```
# Extract the diagonal elements
D <- matrix(1:8,nrow=4)
D
```

```
      [,1] [,2]
[1,]     1     5
[2,]     2     6
[3,]     3     7
[4,]     4     8
```

```
diag(D)
```

```
[1] 1 6
```

¹Can also be used for dataframes (see later)

```
# Calculate eigenvalues and eigenvectors of A
r <- eigen(A)
r
```

```
eigen() decomposition
$values
[1] 5.3722813 -0.3722813

$vectors
      [,1]      [,2]
[1,] -0.5657675 -0.9093767
[2,] -0.8245648  0.4159736
```

```
# Eigenvalues
```

```
r$values

[1] 5.3722813 -0.3722813
```

```
# Matrix with eigenvectors
```

```
r$vectors

      [,1]      [,2]
[1,] -0.5657675 -0.9093767
[2,] -0.8245648  0.4159736
```

```
# Diagonal Matrix (Similarity Transformation)
```

```
solve(r$vectors) %*% A %*% r$vectors
```

```
      [,1]      [,2]
[1,] 5.372281e+00  3.907616e-18
[2,] 1.650150e-16 -3.722813e-01
```

Note that under the hood R calls **BLAS** and **LAPACK**.

```
# Find the version used of BLAS and LAPACK
```

```
La_library()
```

```
[1] "/usr/lib/x86_64-linux-gnu/openblas-pthread/libopenblas-p0.3.20.so"
```

```
extSoftVersion()["BLAS"]
```

```
BLAS
"/usr/lib/x86_64-linux-gnu/openblas-pthread/libblas.so.3"
```

1.3.2 Exercises

- Linear regression:

– Step 1:

Create a **synthetic** data set by executing the following R code:

```
x <- seq(from=0, to=20.0, by=0.25)
a <- 2.0
b <- 1.5
c <- 0.5
y <- a + b*x + c*x^2 + rnorm(length(x))
```

– Step 2:

Our goal is to use the following linear model, i.e.:

$$Y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \epsilon_i$$

or in matrix form:

$$Y = X\beta + \epsilon \quad (1)$$

to fit the previously generated data set.

In Eq.(1), we have:

- * Y : a $n \times 1$ column vector.
- * X : a $n \times 3$ matrix (also known as the **design matrix**)
- * β : a 3×1 column vector.
- * ϵ is : a $n \times 1$ column vector and $\sim N(0, \sigma^2)$

An estimate for β ($\hat{\beta}$) can be found (using Least-Squares, MLE see e.g. (Seber & Lee, 2012)) and has the following form:

$$\hat{\beta} = (X^T X)^{-1} X^T Y \quad (2)$$

where:

the column vector Y is given by:

$$Y := \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

and the matrix X^2 takes the following form:

$$X := \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \vdots & \vdots & \vdots \\ 1 & x_n & x_n^2 \end{bmatrix}$$

Note:

1. The underlying assumption for Eq. (2) is that the inverse of the matrix $(X^T X)$ exists. This is the case iff the $\text{rank}(X^T X) = \text{rank}(X)$ is maximal or if the columns of the matrix X are linearly independent.

2. In the field of machine learning (ML), the vector β is split into 2 parts: the scalar $b := \beta_0$ (**bias**) and the remainder of the vector β , i.e. (w) also known as the **weight** vector.

In the current exercise, the column vector w is given by $(\beta_1 \beta_2)^T$.

Calculate $\hat{\beta}$ using Eq.(2).

An estimate for the residuals ($\hat{\epsilon}$) is given by:

$$\hat{\epsilon} = Y - X\hat{\beta} \quad (3)$$

Calculate $\hat{\epsilon}$ using Eq.(3).

²This is known as a **Vandermonde** matrix.

– Step 3:

You can check your results using the following R code.

```
myquadfit <- lm(y ~ x + I(x^2))
cat(sprintf("The estimates for beta:\n"))
cat(myquadfit$coefficients)
cat(sprintf("The residuals:\n"))
cat(myquadfit$residuals)
```

1.4 Hash tables/dictionaries

We can also use hashes for matrices. We can select one or both dimensions. To create hashes, for: - rows: use **rownames** - columns: use **colnames**

To remove the hash, use the **NULL** (like for vectors).

1.4.1 Examples

```
A1 <- c(0, 5471.52, 5091.57, 5392.82,
        5416.45, 4584.33, 4904.83, 3851.73)
A2 <- c(5471.52, 0, 1315.28, 927.35,
        1505.11, 944.40, 1157.42, 1945.42)
A3 <- c(5091.57, 1315.28, 0, 2166.00,
        2724.01, 1571.76, 293.52, 1240.77)
A4 <- c(5392.82, 927.35, 2166.00, 0,
        577.85, 973.23, 1947.28, 2422.32)
A5 <- c(5416.45, 1505.11, 2724.01, 577.85,
        0, 1366.63, 2490.97, 2838.62)
A6 <- c(4584.33, 944.40, 1571.76, 973.23,
        1366.63, 0, 1290.15, 1474.26)
A7 <- c(4904.83, 1157.42, 293.52, 1947.28,
        2490.97, 1290.15, 0, 1064.41)
A8 <- c(3851.73, 1945.42, 1240.77, 2422.32,
        2838.62, 1474.26, 1064.41, 0)
```

```
dist <- rbind(A1,A2,A3,A4,A5,A6,A7,A8)
dist
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]
A1	0.00	5471.52	5091.57	5392.82	5416.45	4584.33	4904.83	3851.73
A2	5471.52	0.00	1315.28	927.35	1505.11	944.40	1157.42	1945.42
A3	5091.57	1315.28	0.00	2166.00	2724.01	1571.76	293.52	1240.77
A4	5392.82	927.35	2166.00	0.00	577.85	973.23	1947.28	2422.32
A5	5416.45	1505.11	2724.01	577.85	0.00	1366.63	2490.97	2838.62
A6	4584.33	944.40	1571.76	973.23	1366.63	0.00	1290.15	1474.26
A7	4904.83	1157.42	293.52	1947.28	2490.97	1290.15	0.00	1064.41
A8	3851.73	1945.42	1240.77	2422.32	2838.62	1474.26	1064.41	0.00

```
# Adding hashes to both rows and columns
```

```
cities <- c("Anchorage","Atlanta","Austin","Baltimore","Boston", "Chicago", "Dallas","Denver")
```

```
rownames(dist) <- cities
colnames(dist) <- cities
dist
```

	Anchorage	Atlanta	Austin	Baltimore	Boston	Chicago	Dallas	Denver
Anchorage	0.00	5471.52	5091.57	5392.82	5416.45	4584.33	4904.83	3851.73
Atlanta	5471.52	0.00	1315.28	927.35	1505.11	944.40	1157.42	1945.42
Austin	5091.57	1315.28	0.00	2166.00	2724.01	1571.76	293.52	1240.77
Baltimore	5392.82	927.35	2166.00	0.00	577.85	973.23	1947.28	2422.32
Boston	5416.45	1505.11	2724.01	577.85	0.00	1366.63	2490.97	2838.62
Chicago	4584.33	944.40	1571.76	973.23	1366.63	0.00	1290.15	1474.26
Dallas	4904.83	1157.42	293.52	1947.28	2490.97	1290.15	0.00	1064.41
Denver	3851.73	1945.42	1240.77	2422.32	2838.62	1474.26	1064.41	0.00

```
dist["Chicago", "Denver"]
```

```
[1] 1474.26
```

```
dist["Austin", "Boston"]
```

```
[1] 2724.01
```

1.5 Arrays

Say something about arrays.

2 Special Data Types (Factors and Date/Time types)

Every R object has attributes (i.e. properties or metadata).
They can be classified as:

- **intrinsic** properties e.g. `length()`
- **external** properties (to be set by the user)

2.1 Attributes

- can be get/retrieved using `attributes()`.
- can be set:
 - individually using `attr()`
 - in generally using `structure()`
- some attributes can (also) be set/unset with **special** functions:
 - names: `names()`
 - dimension: `dim()`
 - comment : `comment()`
 - time series: `tsp()`
 - factor : `factor()` (see next section)

2.1.1 Examples

- 1 attribute:

```
x <- 1:5  
x
```

```
[1] 1 2 3 4 5
```

```
attr(x, 'prop1') <- "hello"  
attributes(x)
```

```
$prop1  
[1] "hello"
```

```
x
```

```
[1] 1 2 3 4 5
```

```
attr("prop1")  
[1] "hello"
```

```
attr(x, 'prop1') <- NULL  
attributes(x)
```

```
NULL
```

```
x
```

```
[1] 1 2 3 4 5
```

- more than 1 attribute:


```

y <- 1:8
y

[1] 1 2 3 4 5 6 7 8
y <- structure(y, dim=c(2,4), tag="trial")
y

      [,1] [,2] [,3] [,4]
[1,]    1    3    5    7
[2,]    2    4    6    8
attr(,"tag")
[1] "trial"
attributes(y)

$dim
[1] 2 4

$tag
[1] "trial"
typeof(y)

[1] "integer"
class(y)

[1] "matrix" "array"

```

```

# Remove BOTH attributes
y <- structure(y, dim=NULL, tag=NULL)
y

[1] 1 2 3 4 5 6 7 8
attributes(y)

NULL
typeof(y)

[1] "integer"
class(y)

[1] "integer"

```

- `names()`

```

# Set the names attribute
capitals <- c("Salt Lake City", "Carson City", "Boise", "Santa Fe")
names(capitals) <- c("UT", "NV", "ID", "NM")
capitals

```

UT	NV	ID	NM
"Salt Lake City"	"Carson City"	"Boise"	"Santa Fe"

```
attributes(capitals)
```

```
$names
[1] "UT" "NV" "ID" "NM"
```

```
# Remove the names attribute
names(capitals) <- NULL
capitals
```

```
[1] "Salt Lake City" "Carson City" "Boise" "Santa Fe"
```

- `dim()`

```
x <- 1:12
x
```

```
[1] 1 2 3 4 5 6 7 8 9 10 11 12
```

```
typeof(x)
```

```
[1] "integer"
```

```
class(x)
```

```
[1] "integer"
```

```
# Set the dimension attribute
dim(x) <- c(3,4)
x
```

	[,1]	[,2]	[,3]	[,4]
[1,]	1	4	7	10
[2,]	2	5	8	11
[3,]	3	6	9	12

```
typeof(x)
```

```
[1] "integer"
```

```
class(x)
```

```
[1] "matrix" "array"
```

```
# Remove the dimension attribute
```

```
dim(x) <- NULL
```

```
x
```

```
[1] 1 2 3 4 5 6 7 8 9 10 11 12
```

```
typeof(x)
```

```
[1] "integer"
```

```
class(x)
```

```
[1] "integer"
```

- **comment()**

```
x <- structure(1:6, comment="My vector")
```

```
typeof(x)
```

```
[1] "integer"
```

```
class(x)
```

```
[1] "integer"
```

```
comment(x)
```

```
[1] "My vector"
```

2.2 Factor variables (Categorical variables)

- Factor variables (factors, categorical variables) are discrete variables (i.e not continuous). The factors bear labels (**levels**) which are mapped into **integers**.
- Therefore, factors are stored as integer vector with 2 attributes:
 - **class**= “factor”
 - **levels**: a vector with the “labels”.
- By default (**unordered**) the labels are mapped **alphabetically** to the integers. We can **impose** our own **ordering** between integers and labels (levels).
- Useful functions:
 - **levels()** : provides the levels of a factor
 - **table()**: returns the counts of each level
 - **is.factor()**: tests whether a variable is a factor variable
 - **is.ordered()**: tests whether a variable is an ordered factor variable

2.2.1 Examples

- Creation of an **unordered** factor

```
# Creation of an unordered factor
```

```
temp.data <- c("High", "Low", "VeryHigh", "Low", "VeryLow", "Medium",  
              "VeryHigh", "VeryHigh", "Low", "Low", "Medium", "VeryHigh",  
              "VeryHigh", "VeryHigh", "Low", "High", "VeryLow")
```

```
myfac.temp.data <- factor(temp.data)
```

```
myfac.temp.data
```

```
[1] High      Low      VeryHigh Low      VeryLow Medium VeryHigh VeryHigh  
[9] Low      Low      Medium   VeryHigh VeryHigh VeryHigh Low      High
```

```
[17] VeryLow
Levels: High Low Medium VeryHigh VeryLow
```

```
# by default: the levels are stored ALPHABETICALLY (i.e. unordered)
levels(myfac.temp.data)
```

```
[1] "High"      "Low"       "Medium"    "VeryHigh" "VeryLow"
table(myfac.temp.data)
```

```
myfac.temp.data
      High      Low   Medium VeryHigh  VeryLow
        2        5        2         6         2
```

```
is.factor(myfac.temp.data)
```

```
[1] TRUE
```

```
is.ordered(myfac.temp.data)
```

```
[1] FALSE
```

- Creation of an **ordered** factor

```
# Creation of an unordered factor
temp.data <- c("High", "Low", "VeryHigh", "Low", "VeryLow", "Medium",
              "VeryHigh", "VeryHigh", "Low", "Low", "Medium", "VeryHigh",
              "VeryHigh", "VeryHigh", "Low", "High", "VeryLow")
myfac2.temp.data <- factor(temp.data, ordered=TRUE,
                          levels=c("VeryLow", "Low", "Medium", "High", "VeryHigh"))
myfac2.temp.data
```

```
[1] High      Low      VeryHigh Low      VeryLow Medium  VeryHigh VeryHigh
[9] Low      Low      Medium   VeryHigh VeryHigh VeryHigh Low      High
[17] VeryLow
Levels: VeryLow < Low < Medium < High < VeryHigh
```

```
# The ordering is NOW imposed
levels(myfac2.temp.data)
```

```
[1] "VeryLow" "Low"     "Medium"  "High"    "VeryHigh"
table(myfac2.temp.data)
```

```
myfac2.temp.data
  VeryLow      Low   Medium      High VeryHigh
        2        5        2         2         6
```

```
is.factor(myfac2.temp.data)
```

```
[1] TRUE
```

```
is.ordered(myfac2.temp.data)
```

```
[1] TRUE
```

```
# Stripping a factor to the essentials: integer vector
```

```
attributes(myfac2.temp.data)
```

```
$levels
```

```
[1] "VeryLow" "Low"      "Medium"  "High"    "VeryHigh"
```

```
$class
```

```
[1] "ordered" "factor"
```

```
class(myfac2.temp.data) <- NULL
```

```
levels(myfac2.temp.data) <- NULL
```

```
myfac2.temp.data
```

```
[1] 4 2 5 2 1 3 5 5 2 2 3 5 5 5 2 4 1
```

2.3 Dates and times in R.

- **Date** class :
 - represents calendar dates
 - built on top of doubles with class attribute 'Date'
 - 0 : Jan 1, 1970 (**Unix Epoch time**)
 - **as.Date()**: method to cast string to a Date
- **POSIXct** and **POSIXlt** : date and time
 - **POSIXct**: stores date/time values as the #seconds since Jan. 1, 1970
 - **POSIXlt**: stored as **bluelist** with elements for seconds, minutes, hours, day, month, year, etc.
- **lubridate**: a very useful package for dates and times:

2.3.1 Examples

- **Date**

```
today <- Sys.Date()
```

```
today
```

```
[1] "2024-07-22"
```

```
# Attributes of Date
```

```
class(today)
```

```
[1] "Date"
```

```
attributes(today)
```

```
$class  
[1] "Date"
```

```
unclass(today)
```

```
[1] 19926
```

```
d0 <- structure(0, class='Date')  
d0
```

```
[1] "1970-01-01"
```

```
class(d0)
```

```
[1] "Date"
```

```
typeof(d0)
```

```
[1] "double"
```

```
# Convert a string into a Date  
d1 <- as.Date("2022-01-01")  
d1
```

```
[1] "2022-01-01"
```

```
class(d1)
```

```
[1] "Date"
```

```
typeof(d1)
```

```
[1] "double"
```

- **POSIXct**

```
# Convert a string into a POSIXct object  
now_ct <- as.POSIXct("2018-08-01 22:00", tzzone="MST")  
now_ct
```

```
[1] "2018-08-01 22:00:00 MDT"
```

```
attributes(now_ct)
```

```
$class
```

```
[1] "POSIXct" "POSIXt"
```

```
$tzone
```

```
[1] ""
```

```
typeof(now_ct)
```

```
[1] "double"
```

```
# Removal of the attributes
```

```
attr(now_ct, "tzone") <- NULL
```

```
unclass(now_ct)
```

```
[1] 1533182400
```

Bibliography

Seber G.A.F. & Lee A.J. (2012). Linear Regression Analysis. Wiley Series in Probability and Statistics. Wiley.